



JavaOne™

Sun's 2002 Worldwide Java Developer Conference

How NOT To Write A Microbenchmark

Dr. Cliff Click
Senior Staff Engineer
Sun Microsystems



JavaOne™

Sun's 2002 Worldwide Java Developer Conference

How NOT To Write A Microbenchmark or Lies, Damn Lies, and Microbenchmarks

Microbenchmarks are a sharp knife

Microbenchmarks are like a microscope

Magnification is high,
but what the heck are you looking at?

Like "truths, half-truths, and statistics",
microbenchmarks can be **very** misleading



Learning Objectives

- As a result of this presentation, you will be able to:
 - Recognize when a benchmark lies
 - Recognize what a benchmark can tell you (as opposed to what it purports to tell you)
 - Understand how some popular benchmarks are flawed
 - Write a microbenchmark that doesn't lie (to you)



Speaker's Qualifications

- Dr. Click is a Senior Staff Engineer at Sun Microsystems
- Dr. Click wrote his first compiler at age 16 and has been writing:
 - runtime compilers for 15 years, and
 - optimizing compilers for 10 years
- Dr. Click architected the HotSpot™ Server Compiler



Your JVM is Lousy Because it Doesn't Do... "X"

I routinely get handed a microbenchmark and told "HotSpot doesn't do X"

49% chance it doesn't matter,

49% chance they got fooled,

1% chance HotSpot didn't do "Y" but should

1% chance HotSpot didn't do "X" but should



Agenda

- Recent real-word benchmark disaster
- Popular microbenchmarks and their flaws
- When to disbelieve a benchmark
- How to write your own



What is a Microbenchmark?

- Small program
 - Datasets may be large
- All time spent in a few lines of code
- Performance depends on how those few lines are compiled
- Goal: Discover some particular fact
- Remove all other variables



Why Run Microbenchmarks?

- Discover some targeted fact
 - Such as the cost of 'turned off' Asserts, or
 - Will this inline?
 - Will another layer of abstraction hurt performance?
- Fun
 - My JIT is faster than your JIT
 - My Java is faster than your C
- But dangerous!



How HotSpot Works

- HotSpot is a mixed-mode system
- Code first runs interpreted
 - Profiles gathered
- Hot code gets compiled
- Same code "just runs faster" after awhile
- Bail out to interpreter for rare events
 - Never taken before code
 - Class loading, initializing



Example from Magazine Editor

- What do (turned off) Asserts cost?
- Tiny 5-line function r/w's global variable
- Run in a loop 100,000,000 times
- With explicit check – 5 sec
- With Assert (off) – 0.2 sec
- With no test at all – 5 sec
- What did he really measure?



Assert Example

```
static int sval;    // Global variable
static int test_assert(int val) {
    assert (val >= 0) : "should be positive";
    sval = val * 6;
    sval += 3;      // Read/write global
    sval /= 2;
    return val+2;   //
}
static void main(String args[]) {
    int REPEAT = Integer.parseInt(args[0]);
    int v=0;
    for( int i=0; i<REPEAT; i++ )
        v = test_assert(v);
}
```



Assert Example

```
static int sval;    // Global variable
static int test_assert(int val) {
    assert (val >= 0) : "should be positive";
    sval = val * 6;
    sval += 3;      // Read/write global
    sval /= 2;
    return val+2;
}
static void main(String args[]) {
    int REPEAT = Integer.parseInt(args[0]);
    int v=0;
    for( int i=0; i<REPEAT; i++ )
        v = test_assert(v);
}
```



Assert Example

```
static int sval;    // Global variable
static int test_assert(int val) {
    assert (val >= 0) : "should be positive";
    sval = val * 6;
    sval += 3;      // Read/write global
    sval /= 2;
    return val+2;    //
}
static void main(String args[]) {
    int REPEAT = Integer.parseInt(args[0]);
    int v=0;
    for( int i=0; i<REPEAT; i++ )
        v = test_assert(v);
}
```



Assert Example

```
static int sval;    // Global variable
static int test_explicit(int val) {
    if( val < 0 ) throw ...;
    sval = val * 6;
    sval += 3;      // Read/write global
    sval /= 2;
    return val+2;    //
}
static void main(String args[]) {
    int REPEAT = Integer.parseInt(args[0]);
    int v=0;
    for( int i=0; i<REPEAT; i++ )
        v = test_explicit(v);
}
```



Assert Example (cont.)

- No synchronization \Rightarrow hoist static global into register, only do final write
- Small hot static function \Rightarrow inline into loop
- Asserts turned off \Rightarrow no test in code

```
static int sval;    // Global variable
static void main(String args[]) {
    int REPEAT = Integer.parseInt(args[0]);
    int v=0;

    for( int i=0; i<REPEAT; i++ ) {
        sval = (v*6+3)/2;
        v = v+2;
    }
}
```



Assert Example (cont.)

- Small loop \Rightarrow unroll it a lot
 - Need pre-loop to 'align' loop
- Again remove redundant writes to global static

```
static int sval;    // Global variable
static void main(String args[]) {
    int REPEAT = Integer.parseInt(args[0]);
    int v=0;
    // pre-loop goes here...
    for( int i=0; i<REPEAT; i+=16 ) {
        sval = ((v+2*14)*6+3)/2;
        v = v+2*16;
    }
}
```



Benchmark is 'spammed'

- Loop body is basically 'dead'
- Unrolling speeds up by arbitrary factor
- Note: $0.2 \text{ sec} / 100\text{Million iters} * 450 \text{ Mhz clock}$
 \Rightarrow Loop runs too fast: $< 1 \text{ clock/iter}$
- But yet...
- ... different loops run 20x faster
- What's really happening?



A Tale of Four Loops

```
static void main(String args[]) {  
    int REPEAT = Integer.parseInt(args[0]);  
    int v=0;  
    for( int i=0; i<REPEAT; i++ )  
        v = test_no_check(v); // hidden warmup loop: 0.2 sec  
    long time1 = System.currentTimeMillis()  
    for( int i=0; i<REPEAT; i++ )  
        v = test_explicit(v); // Loop 1: 5 sec  
    long time2 = System.currentTimeMillis()  
    System.out.println("explicit check time="+...);  
    for( int i=0; i<REPEAT; i++ )  
        v = test_assert(v); // Loop 2: 0.2 sec  
    long time3 = System.currentTimeMillis()  
    System.out.println("assert time=",...);  
    for( int i=0; i<REPEAT; i++ )  
        v = test_no_check(v); // Loop 3: 5 sec  
    long time4 = System.currentTimeMillis()  
    System.out.println("no check time=",...);  
}
```



On-Stack Replacement

- HotSpot is mixed mode: interp + compiled
- Hidden warmup loop starts interpreted
- Get's OSR'd: generate code for middle of loop
- 'explicit check' loop never yet executed, so...
 - Don't inline test_explicit(), no unroll
 - 'println' causes class loading \Rightarrow drop down to interpreter
- Hidden loop runs fast
- 'explicit check' loop runs slow



On-Stack Replacement (con't)

- Continue 'assert' loop in interpreter
- Get's OSR'd: generate code for middle of loop
- 'no check' loop never yet executed, so...
 - Don't inline `test_no_check()`, no unroll
- Assert loop runs fast
- 'no check' loop runs slow



Dangers of Microbenchmarks

- Looking for cost of Asserts...
- But found OSR Policy Bug (fixed in 1.4.1)
- Why not found before?
 - Only impacts microbenchmarks
 - Not found in any major benchmark
 - Not found in any major customer app
- Note: test_explicit and test_no_check: 5 sec
 - Check is compare & predicted branch
 - Too cheap, hidden in other math



Self-Calibrated Benchmarks

- Want a robust microbenchmark that runs on a wide variety of machines
- How long should loop run?
 - Depends on machine!
- Time a few iterations to get a sense of speed
 - Then time enough runs to reach steady state
- Report iterations/sec
- Basic idea is sound, but devil is in the details



Self-Calibrated Benchmarks

- First timing loop runs interpreted
 - Also pays compile-time cost
- Count needed to run reasonable time is low
- Main timed run used fast compiled code
- Runs too fast to get result above clock jitter
- Divide small count by clock jitter \Rightarrow random result
- Report random score



jBYTEmark Methodology

- Count iterations to run for 10 msec
 - First run takes 104 msec
 - Way too short, always = 1 iteration
- Then use 1st run time to compute iterations needed to run 100msec
 - Way too short, always = 1 iteration
- Then average 5 runs
 - Next 5 runs take 80, 60, 20, 20, 20 msec
 - Clock jitter = 1 msec



Non-constant work per iter

- Non-sync'd String speed test
- Run with & without some synchronization
- Use concat, continuously add Strings
- Benchmark requires copy before add, so...
 - Each iteration takes more work than before
- Self-timed loop has some jitter...
 - 10000 main-loop iters with sync
 - 10030 main-loop iters withOUT sync



Non-constant work per iter (con't)

- Work for Loop1:
 - 10000^2 copies, 10000 concats, 10000 syncs
- Work for Loop2:
 - 10030^2 copies, 10000 concats, no syncs
- Loop2 avoids 10000 syncs
- But pays 600,900 more copies
- Benchmark compares times, BUT
- Extra work swamps sync cost!



Beware 'hidden' Cache Blowout

- Work appears to be linear with time
- More iterations should report same work/sec
- Grow dataset to run longer
- Blow out caches!
- Self-calibrated loop jitter can make your data set fit in L1 cache **or not!**
- Jitter strongly affects apparent work/sec rate



Beware 'hidden' Cache Blowout

- jBYTEmark StringSort with time raised to 1 sec
- Each iteration adds a 4K array
- Needs 1 or 2 iterations to hit 100msec
- Scale by 10 to reach 1 sec
- Uses either 10 or 20 4K arrays
- Working set is either 40K or 80K
- 64K L1 cache



Explicitly Handle GC

- GC pauses occur at unpredictable times
- GC throughput is predictable
- Either don't allocate in your main loops
 - So no GC pauses
- OR run long enough to reach GC steady state
 - Each run spends roughly same time doing GC



CaffeineMark3 Logic Benchmark

- Purports to test logical operation speed
- With unrolling OR constant propagation, whole function is dead (oops!)
- First timing loop: 900K iterations to run 1 sec
- Second loop runs 2.7M iterations in 0.685 sec
- Score reports as negative (overflow) or huge
- Overall score dominated by Logic score



Know what you test

- SpecJVM98 209_db
 - 85% of time spent in shell_sort
- Scimark2 MonteCarlo simulation
 - 80% of time spent in sync'd Random.next
- jBYTEmark StringSort
 - 75% of time spent in byte copy
- CaffeineMark Logic test
 - "He's dead, Jim"



How To Write a Microbenchmark (general – not just Java)

- Pick suitable goals!
 - Can't judge web server performance by jBYTEmarks
 - Only answer very narrow questions
 - Make sure you actually care about the answer!
- Be aware of system load, clock granularity
- Self-calibrated loops always have jitter
 - Sanity check the time on the main run!
- Run to steady state, at least 10 sec



How To Write a Microbenchmark (general – not just Java)

- Fixed size datasets (cache effects)
 - Large datasets have page coloring issues
- Constant amount of work per iteration
 - Or results not comparable
- Avoid the 'eqntott Syndrome'
 - Capture 1st run result (interpreted)
 - Compare to last result



How To Write a Microbenchmark (general – not just Java)

- Avoid 'dead' loops
 - These can have infinite speedups!
 - Print final answer
 - Make computation non-trivial
- Or handle dead loops
 - Report when speedup is unreasonable
 - CaffeineMark3 & JavaGrande Section 1 mix Infinities and reasonable scores



How To Write a Microbenchmark (Java-specific)

- Be explicit about GC
- Thread scheduling is not deterministic
- JIT performance may change over time
- Warmup loop+test code before ANY timing
 - (HotSpot specific)



OSR Bug Avoidance

- Fixed in 1.4.1
- Write 1 loop per method for Microbenchmarks
- Doesn't matter for 'real' apps
 - Don't write 1-loop-per-method code to avoid bug



Summary

- Microbenchmarks can be easy, fun, informative
- They can be very misleading! Beware!
- Pick suitable goals
- Warmup code before timing
- Run reasonable length of time
- Fixed work / fixed datasets per iteration
- Sanity-check final times and results
- Handle 'dead' loops, GC issues



If You Only Remember One Thing...

Put Microtrust in a Microbenchmark



Q&A



JavaOneSM

Sun's 2002 Worldwide Java Developer Conference

BEYOND BOUNDARIES

Session #